

Class 6: R functions

Gaeun: A16814573

Today we are going to explore R functions and begin to think about writing our own functions.

Let's start simple and write our first function to add some numbers.

Every function in R has at least 3 things:

- a **name**, we pick this
- one or more input **arguments**
- the **body**, where the work gets done

```
# y has a default value of 1, z has a default value of 0
add <- function(x, y=1, z=0) {
  x + y + z
}
```

Now let's try it out.

```
add(10, 1)
```

```
[1] 11
```

```
add(x=c(10,1,1,10), y=1)
```

```
[1] 11  2  2 11
```

```
add(10)
```

```
[1] 11
```

```
add(10, 10)
```

```
[1] 20
```

```
add(10, 10, 20)
```

```
[1] 40
```

```
# na.rm overrides the NA  
mean( c(10, 10, NA), na.rm=TRUE)
```

```
[1] 10
```

Lab sheet work

Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped. Your final function should be adequately explained with code comments and be able to work on an example class gradebook such as this one in CSV format: “<https://tinyurl.com/gradeinput>”

```
# Example input vectors to start with  
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)  
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)  
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

Begin by calculating the average for student 1

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
mean(student1)
```

```
[1] 98.75
```

Try on student 2

```
student2
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
mean(student2, na.rm=TRUE)
```

```
[1] 91
```

Try on student 3

```
student3
```

```
[1] 90 NA NA NA NA NA NA NA
```

```
mean(student3, na.rm=TRUE)
```

```
[1] 90
```

Hm.... this sucks! I need to try something else and come back to this issue of missing values (NAs).

We also want to drop the lowest score from a given student's set of scores.

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

```
# this removes the 8th value  
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

We can try the `min()` function to find the lowest score.

```
min(student1)
```

```
[1] 90
```

I want to find the location of the min value, not the value itself. For this, I can use `which.min()`.

```
student1
```

```
[1] 100 100 100 100 100 100 100 100 90
```

```
which.min(student1)
```

```
[1] 8
```

Let's put these two things together.

```
which.min(student1)
```

```
[1] 8
```

```
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

```
mean(student1[-8])
```

```
[1] 100
```

```
# or  
min.ind <- which.min(student1)  
mean(student1[-min.ind])
```

```
[1] 100
```

```
# or  
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Now trying on student 2. But we need to deal with NA (missing values) somehow..

One idea is we make all the NA values zero.

```
x <- student2  
x
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
is.na(x)
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

```
x[is.na(x)]
```

```
[1] NA
```

```
x[!is.na(x)]
```

```
[1] 100 90 90 90 90 97 80
```

So far we have a working snippet.

```
# Find NAs in `x` and makes them 0  
x[is.na(x)] <- 0  
  
# Drops lowest value and finds mean  
mean(x[-which.min(x)])
```

```
[1] 91
```

Now turn it into a function

```
grade <- function(x) {  
  x[is.na(x)] <- 0  
  mean(x[-which.min(x)])  
}
```

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

But instead of calling each student, we want the function to grade all the students at once.

Now `apply()` to our class gradebook

```
gradebook <- read.csv("https://tinyurl.com/gradeinput",  
                      row.names = 1)  
gradebook
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	NA	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	NA
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	NA	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

To use the `apply()` function on this `gradebook` dataset, I need to decide whether I want to “apply” the `grade()` function over the rows (1) or columns (2) of the `gradebook()`.

```
ans <- apply(gradebook, 1, grade)
ans
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Q2. Using your `grade()` function and the supplied `gradebook`, Who is the top scoring student overall in the `gradebook`?

```
which.max(ans)
```

```
student-18
18
```

Q3. From your analysis of the `gradebook`, which homework was toughest on students (i.e. obtained the lowest scores overall)?

First use the `apply` function and change the `margin` to 2 so that it can indicate the columns (homework).

```
hw <- apply(gradebook, 2, grade)
hw
```

hw1	hw2	hw3	hw4	hw5
89.36842	76.63158	81.21053	89.63158	83.42105

The problem is that it doesn't take into account that that NA homeworks are counted as 0.

```
masked_gradebook <- gradebook
masked_gradebook[is.na(masked_gradebook)] = 0
hw <- apply(masked_gradebook, 2, mean)
hw
```

hw1	hw2	hw3	hw4	hw5
89.00	72.80	80.80	85.15	79.25

Now find the homework score that obtained the lowest score.

```
which.min(hw)
```

```
hw2  
2
```

I could modify the `grade()` function to this too - i.e. not drop the lowest options

```
grade2 <- function(x, drop.low=TRUE) {  
  
  # Finds NAs in `x` and makes them 0  
  x[is.na(x)] <- 0  
  
  if(drop.low) {  
    cat("Hello low")  
    # Drop lowest value and find mean  
    out <- mean(x[-which.min(x)])  
  
  } else {  
    out <- mean(x)  
    cat("No low")  
  }  
  return(out)  
}  
  
grade2(student1, FALSE)
```

No low

```
[1] 98.75
```

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)?

The function to calculate correlations in R is called `cor()`

```
x <- c(100, 90, 80, 100)  
y <- c(100, 90, 80, 100)  
z <- c(80, 90, 100, 10)
```



```
# 0 means no correlation at all, 1 means perfectly correlated, -1 means perfectly anti-correlated
cor(x,y)
```

```
[1] 1
```

```
cor(x,z)
```

```
[1] -0.6822423
```

```
cor(ans, masked_gradebook$hw1)
```

```
[1] 0.4250204
```

I want to `apply()` the `cor()` function over the `masked_gradebook()` and use the `ans` scores for the class.

```
predict <- apply(masked_gradebook, 2, cor, y=ans)
```

Find the highest homework correlation.

```
which.max(predict)
```

```
hw5
```

```
5
```