

# mini-project

Gaeun Jun - A16814573

## 1. Exploratory data analysis

```
# Save your input data file into your Project directory
fna.data <- "https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv"
```

```
# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names=1)
```

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438

843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0	0.1622			0.6656
842517	158.80	1956.0	0.1238			0.1866
84300903	152.50	1709.0	0.1444			0.4245
84348301	98.87	567.7	0.2098			0.8663
84358402	152.20	1575.0	0.1374			0.2050
843786	103.40	741.6	0.1791			0.5249
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
84358402	0.4000		0.1625			0.2364
843786	0.5355		0.1741			0.3985
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

Creating an unsupervised data set that doesn't contain the diagnosis.

```
# Using -1 to remove the first column of diagnosis
wisc.data <- wisc.df[, -1]
```

Creating a diagnosis vector that contains the diagnosis column.

```
diagnosis <- wisc.df[,1]
diagnosis <- as.factor(diagnosis)
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in the wisc.data dataset.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

212 of the observations have a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `__mean`?

```
column_indices <- grep("__mean", colnames(wisc.data))
which_column <- colnames(wisc.data)[column_indices]
which_column
```

```
[1] "radius_mean"          "texture_mean"          "perimeter_mean"
[4] "area_mean"            "smoothness_mean"       "compactness_mean"
[7] "concavity_mean"       "concave.points_mean"   "symmetry_mean"
[10] "fractal_dimension_mean"
```

```
length(which_column)
```

```
[1] 10
```

There are 10 variables/features in the data that are suffixed with `__mean`.

## 2. Principal Component Analysis

### Performing PCA

Checking the mean and standard deviations of the wisc.data columns.

```
# Check column means and standard deviations  
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03

compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
# Perform PCA on wisc.data
wisc.data <- scale(wisc.data)
wisc.pr <- prcomp(wisc.data)
```

```
# Look at summary of results
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

0.4427 is captured by PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


$0.4427 + 0.1897 + 0.09393 = 0.726$ . 3 principal components are required to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

$0.4427 + 0.1897 + 0.09393 + 0.06602 + 0.05496 + 0.04025 + 0.02251 = 0.91007$ . 7 principal components are required to describe at least 90% of the original variance in the data.

## Interpreting PCA results

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

Something that stands out to me about this plot is how much clutter there is, making it impossible to read any of the features. It is difficult to understand because there are so many elements in the plot and overlaps that it is not rational to try to comprehend the plot.


Generating a plot that's easier to see.

```
# Scatter plot observations by components 1 and 2
plot(wisc.pr$x[,1], wisc.pr$x[,2], col = diagnosis,
     xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diagnosis,
     xlab = "PC1", ylab = "PC3")
```



Compared to the plot for principal components 1 and 2, this plot seems to have shifted down and slightly more cluttered. This is because PC2 describes more variance in the original data.

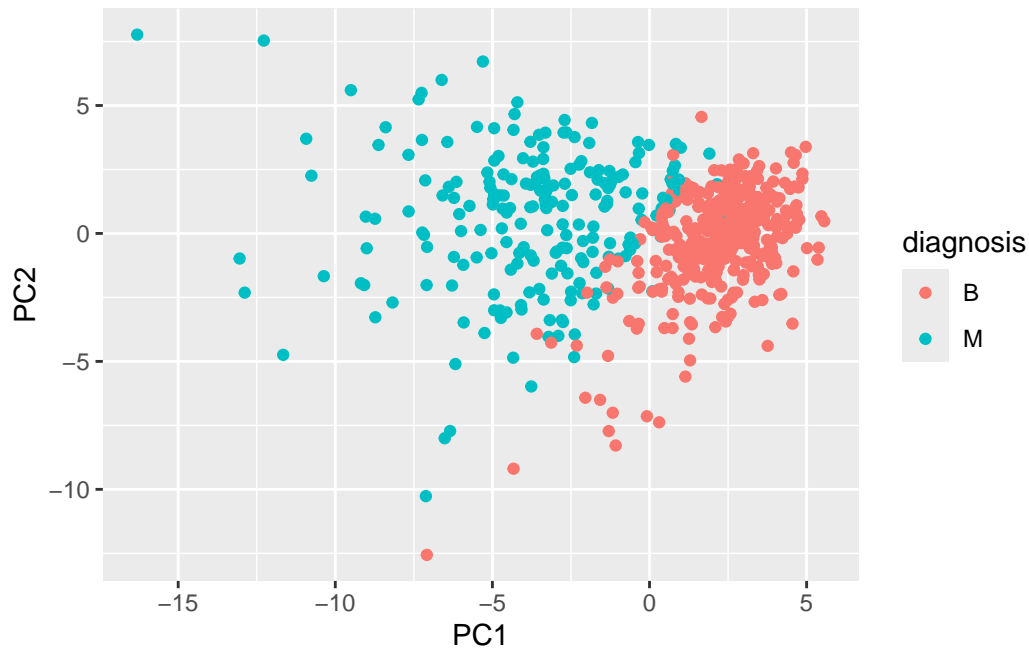
### Using the ggplot2 package

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

library(ggplot2)

ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```






### Variance explained

```
# Calculating the variance of each principal component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```




CRAN packages

```
## ggplot based graph  
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



### Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

```
-0.261
```

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

Using the scree plot, I can calculate how many PC add up to explain at least 80% of the variance of the data:  $44.3 + 19 + 9.4 + 6.6 + 5.5 = 84.8$ . You need at least 5 principal components.

### Hierarchical clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```


```
# Calculate the distances between all pairs of observations
data.dist <- dist(data.scaled)
```

```
# Create a hierarchical clustering model using complete linkage
wisc.hclust <- hclust(data.dist, method = "complete")
```

### Results of hierarchial clustering

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```



At height = 19, the clustering model has 4 clusters

## Selecting number of clusters

Using `cutree()` to cut the tree so that it has 4 clusters.

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
```

```
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.hclust.clusters	B	M	
1	12	165	
2	2	5	
3	343	40	
4	0	2	

```
wisc.hclust.clusters_test <- cutree(wisc.hclust, k=5)  
table(wisc.hclust.clusters_test, diagnosis)
```

	diagnosis		
wisc.hclust.clusters_test	B	M	
1	12	165	
2	0	5	
3	343	40	
4	2	0	
5	0	2	

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

When I cut the amount of clusters to 5 clusters I get a better match because cluster 2 shows a clearer distinction between malignant and benign cases (while the other clusters remain the same) while cluster 2 in the 4 clusters show less of a variance.

## Using different methods

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```

methods <- c("single", "complete", "average", "ward.D2")

results <- list()

# Loop through the methods and store the clusters
for (method in methods) {
  # Create the hierarchical clustering model
  hclust_model <- hclust(data.dist, method = method)

  clusters <- cutree(hclust_model, k = 4)

  results[[method]] <- table(clusters, diagnosis)
}

results

```

```

$single
      diagnosis
clusters  B   M
1 356 209
2    1    0
3    0    2
4    0    1

```

```

$complete
      diagnosis
clusters  B   M
1   12 165
2    2    5
3 343  40
4    0    2

```

```

$average
      diagnosis
clusters  B   M
1 355 209
2    2    0
3    0    1
4    0    2

```

```

$ward.D2
      diagnosis

```

clusters	B	M
1	0	115
2	6	48
3	337	48
4	14	1

Out of these methods, the “complete” and “ward.D2” methods give me the best results because they show the clearest distinction between the benign and malignant cases. However, it’s hard to compare these two because while one might have a bigger discrepancy between the two, the other has a clearer indication that a certain cluster corresponds to either malignant or benign cells.

## OPTIONAL: K-means clustering

### K-means clustering and comparing results

Creating a k-means model on wisc.data

```
scaled_data <- scale(wisc.data)
wisc.km <- kmeans(scaled_data, centers= 2, nstart= 20)
```

```
table(wisc.km$cluster, diagnosis)
```

	diagnosis	
	B	M
1	14	175
2	343	37

Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results?

```
wisc.hclust.clusters.2 <- cutree(wisc.hclust, k=2)
table(wisc.hclust.clusters.2, diagnosis)
```

	diagnosis	
wisc.hclust.clusters.2	B	M
1	357	210
2	0	2



Compared to the hclust results, k-means definitely separates the two diagnoses a lot better, because the first cluster clearly shows an indication towards benign cells and cluster 2 showing an indication towards malignant cells. In the hclust results, when I use 2 clusters, the first cluster is pretty divided between the two cells.

```
table(wisc.hclust.clusters, wisc.km$cluster)
```

```
wisc.hclust.clusters  1  2
                     1 160 17
                     2   7  0
                     3  20 363
                     4   2  0
```

## 5. Combining methods

### Clustering on PCA results

```
pca_result <- prcomp(scaled_data, center = TRUE, scale. = TRUE)


pve <- summary(pca_result)$importance[3, ]
components_needed <- which(pve >= 0.90)[1]

pca_data <- pca_result$x[, 1:components_needed]

wisc.pr.hclust <- hclust(dist(pca_data), method = "ward.D2")

plot(wisc.pr.hclust)
```

## Cluster Dendrogram



```
dist(pca_data)
hclust (*, "ward.D2")
```


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1  2
216 353
```


```
table(grps, diagnosis)
```

```
      diagnosis
grps   B    M
  1  28 188
  2 329  24
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



Swapping the colors of the clusters, such that cluster 2 comes first


```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



```
## Use the distance along the first 7 PCs for clustering
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method="ward.D2")
```

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

```
table(wisc.pr.hclust.clusters, diagnosis)
```

```
      diagnosis
wisc.pr.hclust.clusters  B   M
1      28 188
2     329  24
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

It separates out the four clusters as the same as k-means. It separates it out pretty well as the first cluster indicates malignant cells and the second cluster indicates benign cells.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.km$cluster, diagnosis)
```

```
      diagnosis
      B    M
1     14 175
2    343  37
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
      diagnosis
wisc.hclust.clusters  B    M
1      12 165
2       2   5
3     343  40
4       0   2
```

Both k-means and hierarchical clustering can separate the diagnosis well, because each cluster tends to lean towards a diagnosis with a good proportion. Even `ward.d2` can separate the diagnosis well.

## 6. Sensitivity/Specificity

Sensitivity: cluster with predominantly malignant cells / total number of known malignant samples

```
sum(diagnosis == "M")
```

```
[1] 212
```

The different methods/procedures for sensitivity:

hcluster:  $165/212 = 0.778$

k-means:  $175/212 = 0.825$

K-means results in a clustering model with the best sensitivity.

Specificity: cluster with predominantly benign cells / total number of known benign samples

```
sum(diagnosis == "B")
```

```
[1] 357
```

hcluster:  $343/357 = .961$

k-means:  $343/357 = .961$


Both hcluster and k-means result in a clustering model with the best specificity.

## 7. Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	-10.76452	-10.093978	-0.5897994	-4.164748	10.61922	-1.630738	0.03566861
[2,]	-18.09606	-9.967098	-2.1549431	-4.006848	6.69687	-2.034714	1.25088149
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	0.7308658	-1.580861	3.166451	-0.7167150	3.850569	-0.8259764	1.0195729
[2,]	0.6308585	-1.155629	3.608207	-0.3405375	2.288732	-0.3976672	0.1347203
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
[1,]	3.735687	-4.068783	1.0877034	0.9985959	1.022760	-2.430215	-1.295749
[2,]	3.543905	-3.749616	0.7613603	1.1763217	1.366702	-2.609643	-1.541050
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
[1,]	-1.348026	-0.7388274	-1.083000	-0.4220831	-1.892993	-1.176056	0.05527974
[2,]	-1.424290	-0.7591376	-1.439202	-0.6508838	-1.981711	-1.397390	0.18112357
	PC29	PC30					
[1,]	0.2658028	0.05162840					
[2,]	0.2842191	0.02734355					

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

Patient 2 should prioritize for follow-up because they have more malignant cells while patient 1 has benign cells.